

CONTENTS

Contents	1
1 Security	2
1.1 Scope and Goals	2
1.2 Scope	2
1.3 Goals	3
2 Background	3
2.1 Trusted Computing	3

the problem of how multiple stakeholders could cooperate to perform collaborative computation on remote computers (6) and how they could trade-off between security and performance (7). Specifically, how they might agree on which security mechanisms they want to rely on to protect their workloads while gaining access to hardware accelerators. Finally, I deal with practical issues that limit the usage of integrity enforcement and monitoring techniques in production. First, I investigate how to safely install software updates on an integrity-enforced operating system (8), i.e., I look for a solution in which a remote verifier who monitors the integrity of the operating system can ensure that the new integrity state is a result of the trusted update and not of an attack. Second, I check how a system owner could in practice manage a group of resources, i.e., define, configure, and monitor remote computers that differ in terms of running workloads and applied security mechanisms (9).

1.3 Goals

The main goal of this thesis is to build a framework to harden high-assurance security systems. The design goals are:

- **Security.** The framework should provide strong security guarantees to high-assurance security systems. It should allow individual processes to be run in isolation from privileged software (under the trusted execution environment threat model) and on top of a trustworthy operating system (under the trusted computing threat model).
- **Attestation.** The system owner should obtain technical assurance that the high-assurance security systems execute in well-defined geographic locations inside an execution environment meeting his security requirements.
- **Practicality.** The framework must support running legacy systems without requiring source code changes. It is acceptable to instrument source code at the compilation level or run inside virtual machines. It must also support software updates and incur acceptable, low ($\leq 10\%$) performance overhead.
- **Usability.** The framework must be configurable to individual use cases by allowing users to declaratively state their trust boundaries and make a trade-off between security and performance. It should permit central management (configuration distribution and notification collection) of multiple computing resources.

2 BACKGROUND

High-assurance security systems are deployed as multiple services distributed across multiple computers to ensure high availability, fault tolerance, and resource scalability. As such, their architects, owners, and security officers face the problem of secure remote computation, i.e., how to ensure the confidentiality and integrity of data and code executing on a remote computer? This chapter explores existing techniques that allow users to establish trust with a remote computer, attest to the integrity of software running on such a computer, and protect the integrity and confidentiality of individual applications' code and data against malicious operating systems and administrators. Figure(2) shows existing defense mechanisms used to protect computing systems at different levels.

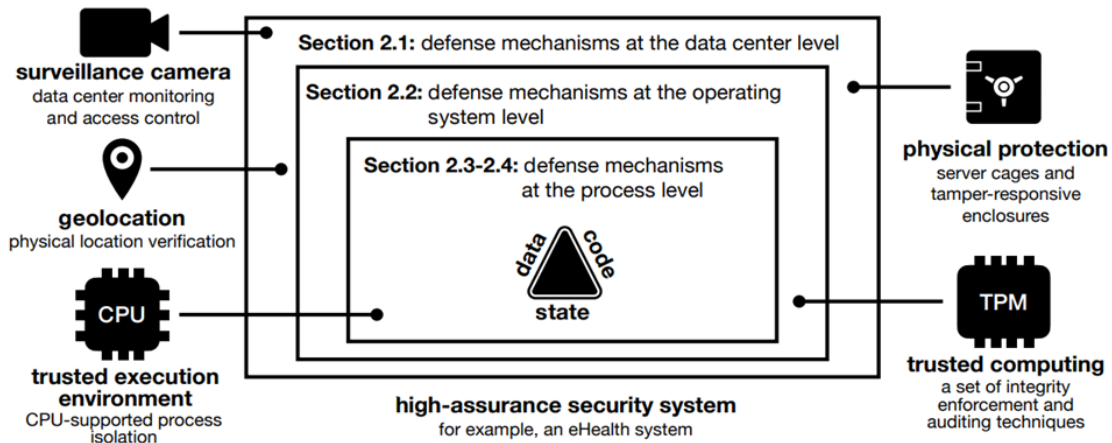


Fig. 2. Overview of mechanisms protecting high-assurance security systems against a wide range of threats

2.1 Trusted Computing

Overview of mechanisms protecting high-assurance security systems against a wide range of threats

- (1) remote attestation, i.e., auditing, of what software has executed on the computer
- (2) integrity enforcement mechanisms ensuring that only expected software in the expected configuration can execute on the computer.

2.1.1 Security Guarantees. TCTs define how to measure, store, enforce, and report the load-time integrity of firmware and software that has been loaded to the computer's memory since the moment a computer was powered-up. The reporting capability (also referred to as auditing or remote attestation) allows verification that the operating system is in the expected, well-defined state, while the enforcement capability prevents the operating system from moving into an untrusted state by refusing to load an unknown, potentially malicious software to the memory. Crucially, the reporting capability verifies that the enforcement mechanism is enabled and certifies this to a remote entity with the help of the secure element.

Table 1. The VM boot time depending on the TPM.

	MC	TPM	IMA	Boot time
No TPM	-	-	-	9.7 sec ($\sigma = 0.1$ sec)
swTPM	-	+	+	14.0 sec ($\sigma = 0.1$ sec)
TRIGLAV				
No MC	-	+	+	14.1 sec ($\sigma = 0.1$ sec)
With MC	+	+	+	50.8 sec ($\sigma = 0.1$ sec)
Fast MC	+	+	+	15.8 sec (estimate)

Table(1) shows how TRIGLAV impacts VM boot times. As a reference, we measure the boot time of a VM without any TPM attached. Then, we run experiments in which a VM has access to different implementations of a software-based TPMs. Except for the reference measurement, the Linux IMA is always turned on. Each VM has access to all available cores and 4 GB of memory. As the guest operating system, we run Ubuntu 18.10, a Linux distribution with a pre-installed tool (systemd-analyze) to calculate system boot times. Shannon's entropy formula:

$$H(X) = - \sum p(x) * \log_2 p(x),$$

$H(X)$ is the entropy; $p(x)$ is the probability of symbol x .

$$hash = (S[i] + d * (hash - S[i - len + 1])) \mod q,$$

where S is the string; d is the alphabet size; q is a prime number.